

PIPE-RDF: Execution-Grounded Generation of Schema-Specific NL-SPARQL Benchmarks

Anonymous ACL submission

Abstract

Natural-language access to RDF knowledge graphs depends on evaluation sets whose queries actually run on the graph being tested. Existing KGQA benchmarks provide useful shared tasks, but their schemas, namespaces, predicates, and query distributions often differ from the graphs used in practice. We introduce PIPE-RDF, an execution-grounded workflow for building schema-specific natural-language/SPARQL benchmarks from a target RDF graph. PIPE-RDF starts with reverse queries that populate deterministic binding banks, uses category-aware retrieval to supply schema-matched examples, and applies controlled LLM generation to produce candidate question-query pairs. Each candidate is accepted only after passing predicate and type checks, deduplication, parsing, execution, answer-shape validation, and non-empty-result checks. Across a compact company-location schema and a 25M-triple LDBC Semantic Publishing Benchmark graph, PIPE-RDF produces 3,600 balanced pairs across nine query categories with no parse, execution, or empty-answer failures in the released artifacts. Cross-model probes, a dual LLM-judge audit, and downstream prompting experiments indicate that the resulting benchmarks are robust, semantically aligned, and useful as schema-matched examples for NL-SPARQL evaluation. Anonymous code is available at <https://anonymous.4open.science/r/PIPE-RDF-E5BD/README.md>.

1 Introduction

RDF knowledge graphs expose structured data through SPARQL, but natural-language access to those graphs remains highly schema-dependent. A natural-language question can be fluent while its corresponding query uses the wrong predicate, assumes a missing type, or returns an answer set that cannot exist in the target graph. This makes benchmark construction a first-order problem for

NL-SPARQL work: an evaluation set must reflect the target ontology, prefixes, available attributes, and intended query categories before it can be used to compare parsers or retrieval strategies.

Public KGQA benchmarks such as LC-QuAD, QALD-9-plus, and Mintaka provide valuable shared test beds, but they are tied to fixed public schemas and entity distributions (Trivedi et al., 2017; Perevalov et al., 2022; Sen et al., 2022). They do not directly answer a common methodological question: given a new RDF schema, how can we construct a balanced, executable, and auditable NL-SPARQL benchmark without manually authoring every query? Modern large language models (LLMs) can help generate candidate questions and queries (Brown et al., 2020; Touvron et al., 2023), but unconstrained generation is prone to unsupported entities, invalid predicates, and semantically plausible but unanswerable queries.

PIPE-RDF addresses this benchmark-construction problem with a pipeline that grounds generation in graph facts. The central mechanism is reverse querying: before a question-query pair is accepted, the system first finds entity bindings that satisfy the intended template in the target graph. Category-balanced generation, retrieval-augmented prompting, predicate/type whitelists, deduplication, execution validation, and repair are then used to produce benchmark artifacts that are not merely grammatical, but executable against a fixed RDF endpoint.

Contributions. We:

- propose a three-phase pipeline for schema-specific NL-SPARQL benchmark generation with reverse-query grounding and execution validation;
- evaluate PIPE-RDF on a fixed company-location schema and a large LDBC Semantic Publishing Benchmark (SPB) graph with

a broader schema profile, producing 3,600 accepted Phase-3 pairs across nine categories;

- report operational metrics including parse validity, execution success, empty-result rates, repair rates, materialized generation attempts, endpoint latency, and predicate usage;
- add cross-model robustness probes and a compact downstream utility evaluation for zero-shot versus retrieval-augmented NL-to-SPARQL prompting;
- report a dual LLM-judge semantic audit with agreement, Cohen’s κ , consensus error categories, and the exact judging prompt.

2 Background and Related Work

Benchmark coverage and query taxonomies. NL-SPARQL and KGQA benchmarks have been developed through several complementary traditions. LC-QuAD, LC-QuAD 2.0, QALD-9-plus, and Spider4SPARQL provide question-query pairs over public graphs and established ontologies (Trivedi et al., 2017; Dubey et al., 2019; Perevalov et al., 2022; Kosten et al., 2023). CFQ uses SPARQL over Freebase-style data to stress compositional generalization (Keysers et al., 2020), and KQA Pro provides explicit compositional programs plus SPARQL for complex KBQA (Cao et al., 2022). Mintaka explicitly balances query categories (counting, comparative, superlative, ordinal, multi-hop, intersection, difference/negation, yes/no, and generic) to avoid skewed evaluation (Sen et al., 2022). These datasets are valuable shared test beds, but their schemas and entity distributions are fixed once released. PIPE-RDF instead targets the complementary problem of constructing an executable benchmark for a user-specified RDF schema.

Template and workload generation. Earlier SPARQL dataset and workload generators show that template instantiation and workload design are central for RDF evaluation. Neural SPARQL Machines generate synthetic NL-SPARQL training pairs from templates over DBpedia-style data (Soru et al., 2017). RDF engine benchmarks such as SP2Bench, LDBC SPB, and FEASIBLE study scalable SPARQL workload generation and query-log or feature-driven benchmark construction (Schmidt et al., 2008; Linked Data Benchmark Council, 2026; Saleem et al., 2015). PIPE-RDF

builds on this lineage but adds LLM-era NL generation, category-specific retrieval, binding-bank construction, and strict accepted-record validation for schema-specific KGQA artifacts rather than triplestore throughput alone.

Retrieval augmentation and prompting.

Retrieval-augmented generation (RAG) improves text-to-query performance by conditioning on relevant examples and schema cues (Lewis et al., 2021). In practice, RAG offers a low-cost alternative to fine-tuning when schema-specific examples are scarce. PIPE-RDF maintains category-specific retrieval banks to provide structurally similar NL-SPARQL pairs during generation, reducing cross-category drift and improving template grounding.

Validity and evaluation. Grammar-constrained decoding (e.g., PICARD) improves syntactic correctness for structured query generation (Scholak et al., 2021). SPARQL-specific constrained decoding such as SPARKLE integrates KG structure during inference to reduce inexecutable query generation (Lee and Shin, 2025). SPARQL-specific metrics such as SP-BLEU/SP-F1 normalize variables for fairer matching (Rony et al., 2022), while code-oriented structural metrics such as CodeBLEU inspire future extensions (Ren et al., 2020). LLM-as-judge protocols such as G-Eval and MT-Bench show that strong model judges can provide scalable rubric-based judgments when prompts and outputs are logged (Liu et al., 2023; Zheng et al., 2023). PIPE-RDF combines execution, parsing, structural, and dual-judge semantic metrics so that accepted artifacts can be audited beyond surface validity.

Agentic and verification-heavy NL-SPARQL systems. Recent systems such as GRASP, FIRES-PARQL, and Chatty-KG use LLM agents, decomposition, verification, or conversational orchestration to answer questions over knowledge graphs (Walter and Bast, 2025; Pan et al., 2025; Omar et al., 2025). PIPE-RDF is complementary: it does not propose a new inference-time question-answering agent, but constructs schema-specific, executable benchmark artifacts that can evaluate or supply retrieval/fine-tuning examples for such systems.

Synthetic text-to-query benchmarks. Synthetic benchmark construction is also active outside RDF. For example, SM3-Text-to-Query builds a synthetic multi-model medical text-to-query benchmark across heterogeneous data models (Fürst et al., 2024). PIPE-RDF makes a different trade-off: it targets RDF/SPARQL specifically and prioritizes

Feature	Public KGQA	CFQ/KQA	RDF workloads	PIPE-RDF
Released NL-SPARQL pairs	✓	✓	–	✓
New-schema configuration	–	–/partial	✓	✓
Graph-grounded bindings	partial	partial	✓	✓
Category-balanced generation	partial	✓	–	✓
LLM generation/RAG controls	–	–	–	✓
Execution + semantic audit	partial	partial	✓	✓

Table 1: Comparison of PIPE-RDF with existing KGQA datasets and RDF workload-generation approaches. PIPE-RDF targets schema-specific NL-SPARQL benchmark construction with both graph-grounded generation and accepted-record auditing. “Partial” denotes support that is limited to fixed released benchmarks or does not directly handle user-provided schemas.

graph-grounded executability, schema whitelists, category balance, and released audit logs over unconstrained linguistic breadth.

LLM-driven workload generation for analytics. Recent work on cross-model analytics workload generation uses a two-stage design (reverse query generation followed by NL query generation), explicit schema/strategy controls, and detailed error analyses across model families (Zheng et al., 2024). We adopt the same design principle of schema-grounded workload construction, and the released artifacts log strategy tags and failure modes to make benchmark generation auditable beyond aggregate success rates.

Comparison with existing approaches. Table 1 contrasts PIPE-RDF with prominent KGQA and SPARQL workload-generation methodologies. Our novelty is not a new semantic parser or the isolated use of templates. The contribution is an integrated, reproducible construction pipeline that combines graph-grounded binding banks, category-aware LLM prompting, strict acceptance filters, LLM semantic judging, and run-level observability for new RDF schemas.

3 Benchmark Design Goals

PIPE-RDF is designed to meet four goals for schema-specific benchmark construction: (i) **schema fidelity** - generated queries must be grounded in a fixed schema and valid predicates; (ii) **category balance** - benchmarks should cover complex query types rather than skew toward easy factoids; (iii) **operational observability** - generation, parsing, and execution outcomes must be logged for debugging and cost control; and (iv) **maintainability** - the pipeline should be repeatable as schemas evolve, minimizing manual curation. These goals shape the pipeline phases, validation logic, and artifact outputs.

Category	SPARQL construct	Example pattern
Generic	single triple	Who is a key person at company?
Counting	COUNT	How many companies are in location?
Comparative	boolean comparison	Do company1 and company2 differ in size?
Superlative	ordering	Which company has most employees?
Ordinal	ranked ordering	Which company was founded second earliest?
Multi-hop	join chains	Which location has companies with key persons?
Intersection	conjunction	Which companies are in location and industry?
Difference	boolean exclusion	Are company1 and company2 in different industries?
Yes/No	boolean check	Is company located in location?

Table 2: Category taxonomy used for balanced benchmark generation.

4 Category Taxonomy

To prevent evaluation bias, PIPE-RDF follows a nine-category taxonomy derived from complex KGQA datasets (Sen et al., 2022). Each category maps to distinct SPARQL constructs (e.g., aggregations, ordering, multi-hop joins) and therefore stresses different failure modes. Table 2 summarizes the taxonomy and example question patterns used in Schema C.

5 PIPE-RDF Methodology

PIPE-RDF is a three-phase pipeline (Figure 1) designed to generate schema-specific benchmarks while minimizing unanswerable or out-of-schema queries.

Phase 1: seed generation. The pipeline prompts an LLM to produce NL templates from a schema summary. We represent each template as a tuple $(c, q_\theta, s, r_\theta, v_c)$, where c is the category, q_θ is an NL pattern with typed slots, s is the slot-type map, r_θ is a reverse SPARQL pattern that returns candidate bindings for the slots, and v_c is the category validator. Each template is validated by *reverse querying*: the system generates or renders a SPARQL query that searches the KG for entity bindings satisfying the template. Valid bindings instantiate the template and the resulting NL-SPARQL pairs are stored as seeds for later retrieval.

Phase 2: category-wise seeding. Templates are generated per category and paired with retrieval-augmented prompts using Phase-1 seeds. Outputs are validated by execution, optionally repaired, and stored in category-specific retrieval banks. This step ensures early coverage of complex query structures (e.g., aggregations, comparisons, multi-hop joins) and prevents domination by simple tem-

plates.

Phase 3: full dataset generation. The pipeline generates batches per category using category-specific retrieval banks. Each candidate is parsed, executed, and logged. Records are retained only after strict acceptance checks, then exported as balanced CSV/JSONL artifacts with run summaries.

Prompt construction and retrieval. Prompts combine (i) a task description, (ii) the schema summary, (iii) a small set of retrieved NL-SPARQL examples, and (iv) the target question. Retrieval is category-aware, which narrows the space of candidate structures and reduces cross-category drift. The ARR runs use top- $k = 3$ to balance example coverage against prompt length.

Reverse querying and schema control. Reverse queries build deterministic binding banks for each template before generation. Each bank stores a bounded pool of valid entity/literal bindings, and Phase 3 samples from those banks client-side rather than repeatedly issuing randomized SPARQL queries. This avoids broad ORDER BY RAND() scans on large repositories and makes GraphDB work an explicit upfront stage. We also batch label and type resolution for candidate URIs with VALUES clauses, reducing thousands of small endpoint calls. We enforce explicit predicate/type whitelists and slot-type hints (company, location, person, industry) to ensure every generated query conforms to the target schema. Duplicate questions are filtered using token-level Jaccard similarity (threshold 0.97 in the ARR runs), and a stall limit prevents repeated template regeneration when no new valid rows are found.

Validation and repair. Each generated SPARQL query is parsed and executed; failures trigger a repair prompt that attempts to correct syntax and predicate usage. Repairs are logged and included in run summaries for operational tracking. The expanded runs use a stricter acceptance gate than the original baseline: accepted records must parse, execute, return a non-empty answer object, and match the expected answer type for the category. Boolean ASK queries therefore produce explicit true/false answers rather than empty-result negatives. The primary benchmark is an answerable evaluation split, so zero-count counting questions are excluded from this split rather than mixed with positive counting examples; controlled zero-count diagnostics can be generated as a separate negative split using the same template interface. Post-generation, category-specific pat-




Figure 1: PIPE-RDF workload-construction pipeline: Phase 1 uses reverse query generation to build verified seeds; Phase 2 constructs category-specific seed banks; Phase 3 generates the balanced benchmark with execution-driven validation/repair.

tern enforcement can normalize output forms (e.g., ASK for boolean checks, DISTINCT for set-returning categories, COUNT(DISTINCT ...) for counting queries, and deterministic LIMIT 1 tie-breaking for superlative/ordinal queries). We also run a deterministic same/different lexical-shape check for comparative and difference questions, preventing records whose question asks for sameness while the query tests inequality, or vice versa. While PIPE-RDF does not require constrained decoding, it can integrate grammar-based constraints (e.g., PICARD) or KG-aware decoding (e.g., SPARKLE) to improve parse validity at inference time (Scholak et al., 2021; Lee and Shin, 2025).

Curation controls. The pipeline supports optional curation after each phase, with the ability to discard, edit, or inject examples. This is useful for sparse relations, sensitive attributes, or domain-specific naming conventions where automated validation alone may not capture all semantic constraints.

6 Schemas and Implementation

Fixed schema slice. We derive a company-location mini-slice (Schema C) with 5,000 companies, three explicit classes (dbo:Company, gn:Feature for locations, and foaf:Person for key people), and industry IRIs used as labeled attribute values rather than an explicit typed class. The allowed predicates are dbo:location, dbo:industry, dbo:keyPerson, dbo:foundingYear, dbo:numberOfEmployees, and label predicates (rdfs:label, spb:prefLabel, foaf:name). This fixed schema avoids prompt drift and simplifies validation.

Slice extraction. The mini-slice is sampled from a public RDF source by selecting companies with at least one of the core predicates (location, in-

dustry, key person, founding year, or employee count). Linked entities (locations and key people) are included with minimal labels to support question instantiation without introducing extraneous schema elements. We intentionally exclude unrelated predicates to reduce template drift and ensure that reverse querying remains tractable.

SPB-full schema. To test the same method beyond the compact slice, we also run on a full LDBC Semantic Publishing Benchmark graph with roughly 25M triples. The graph has a much broader profile than Schema C (5,695 distinct predicates and 19,033 distinct types in the profiler output); its top predicates include `rdf:type`, `spb:prefLabel`, `foaf:name`, `dbo:birthPlace`, `dbo:location`, `dbo:foundingYear`, `dbo:industry`, `dbo:keyPerson`, and `dbo:numberOfEmployees`. For comparability, the benchmark-generation task still targets the company/location/person-style query categories, while the schema summary and whitelist are derived from the larger graph. This tests graph scale and broader schema-profile handling while holding the task taxonomy fixed; it is not intended as exhaustive coverage of every predicate or advanced SPARQL construct in SPB.

Schema profiling and curation effort. Schema profiles are produced automatically from the endpoint: `scripts/profile_schema.py` exports predicate counts, type counts, a class–predicate matrix, and a compact schema summary. The remaining schema-specific input is a YAML configuration containing preferred types, slot-type hints, endpoint settings, and whitelist controls. In the ARR runs, Schema C explicitly lists 9 allowed predicates and 3 allowed types; SPB-full uses profiler-derived predicate/type CSVs with top- k limits plus 5 slot-type hints. The hand-authored schema-specific portion is therefore inspectable in the released YAMLs: types, slot hints, and whitelist settings, while predicate/type counts and summaries are generated by the profiler. This keeps portability work localized to the config and profiler outputs rather than hard-coding schema logic in the generator.

Pipeline configuration. The primary ARR runs use 8 Phase-1 templates per category, 10 Phase-1 seed instantiations per template, 50 Phase-2 seeds per category, and 200 accepted Phase-3 targets per category on each schema (3,600 accepted Phase-3 pairs total). Generated SPARQL is capped at 5 results to limit execution cost. Binding banks use bounded deterministic pools and client-side shuffling for diversity.

Class	Key predicates/attributes
<code>dbo:Company</code>	<code>dbo:location</code> , <code>dbo:industry</code> , <code>dbo:keyPerson</code> , <code>dbo:foundingYear</code> , <code>dbo:numberOfEmployees</code> , <code>rdfs:label</code>
<code>gn:Feature</code>	<code>rdfs:label</code> , <code>spb:prefLabel</code>
<code>foaf:Person</code>	<code>foaf:name</code> , <code>rdfs:label</code>

Table 3: Schema C mini-slice explicit classes and allowed predicates. Industry is represented as an IRI-valued company attribute and may not have an explicit `rdf:type`.

Parameter	Value
Phase 1 templates / category	8
Phase 1 seeds / template	10
Phase 2 seeds / category	50
Phase 3 accepted targets / category	200 primary / 50 probe
Binding-bank target rows	500 (Schema C) / 1,500 (SPB-full)
Retrieval top- k	3
SPARQL result cap	5
Dedup Jaccard threshold	0.97

Table 4: Key primary-run configuration. The expanded runs keep the same category taxonomy while scaling accepted Phase-3 targets and moving reverse-query work into deterministic binding-bank construction.

Runtime environment. We run GraphDB 10.6.3 and execute SPARQL with inference disabled. The expanded generation uses Qwen 3.5 models served through vLLM and local BAAI/bge-m3 embeddings. The primary runs use Qwen/Qwen3.5-4B; robustness probes use Qwen/Qwen3.5-2B, Qwen/Qwen3.5-4B, and Qwen/Qwen3.5-9B. All runs are logged with execution latency, parse validity, and answer counts; artifacts are exported as JSONL/CSV and summarized in run-level reports. A compact reproducibility checklist is provided in Appendix Section A.

SPARQL formalism and schema constraints. To ensure semantic correctness, we enforce canonical query patterns:

- **Counting:** Use `COUNT(DISTINCT ?var)` to avoid overcounting due to multiple bindings.
- **Superlative/Ordinal:** Include deterministic tie-breaking via `ORDER BY ... ?entity` to ensure reproducible results.
- **Numeric comparisons:** Cast literals to `xsd:integer` where applicable (e.g., `dbo:numberOfEmployees`).
- **Yes/No:** Use `ASK { }` for explicit boolean returns, avoiding empty-result ambiguity.

Schema constraints are enforced via predicate whitelists: `dbo:location` ranges over `gn:Feature`, `dbo:keyPerson` ranges over `foaf:Person`, and `dbo:foundingYear/dbo:numberOfEmployees` are typed literals. Label resolution prioritizes `foaf:name` for persons, `rdfs:label` for companies, and `spb:prefLabel` for locations.

7 Evaluation Protocol

PIPE-RDF logs operational and structural metrics for generated and accepted records. We track: (i) parse validity and execution success, (ii) non-empty answer coverage, (iii) repair attempts, (iv) materialized generation attempts and failures from run logs, (v) answer counts, (vi) prompt and question lengths, (vii) model and endpoint latency when generated by the active run, and (viii) SPARQL structural complexity (triple pattern counts, FILTER clauses, aggregations). We also compute natural-language surface diagnostics over accepted artifacts: unique-question rate, mean question length, and Distinct-1/Distinct-2. These metrics support both data quality analysis and model evaluation. While PIPE-RDF focuses on execution-driven correctness, the artifacts can be augmented with SPARQL-specific metrics (SP-F1) and structural metrics (CodeBLEU) for model evaluation (Rony et al., 2022; Ren et al., 2020). We also record retrieval scores to analyze the relationship between example similarity and downstream execution success.

Following recent workload-generation evaluation practice (Zheng et al., 2024), the released logs also include strategy tags inferred from query operators (JOIN, FILTER, COUNT, ORDER, NEGATION, ASK, RAG-context). These diagnostics are intended for run auditing and failure analysis rather than as the central acceptance criterion.

Semantic correctness audit. To measure whether generated NL-SPARQL pairs are semantically aligned (not just syntactically valid), we construct a 216-pair stratified audit packet: 2 schemas $\times$ 9 categories $\times$ 12 pairs. We run two independent LLM judges, GPT-5-mini and Grok-4.20-0309 non-reasoning, with the same rubric and temperature 0. Each judge checks whether (i) question intent matches the SPARQL semantics, (ii) entity bindings in the question correspond to VALUES clauses or equivalent query constraints, (iii) the answer type matches the expected return, and (iv)

Schema	Pairs	Cat.	Parse fail	Exec fail	Empty	Repairs
Schema C	1,800	200/200	0	0	0	0
SPB-full	1,800	200/200	0	0	0	0

Table 5: Primary expanded Phase-3 results. Both schemas are balanced at 200 examples per category across nine categories and pass strict parse, execution, answer-type, and category-form checks.

Schema	Runs	Wall h	P3 attempts	Final	Yield%	Repairs (P3)	SPARQL ms
Schema C	3	8.01	2,692	1,800	66.9	32 (1)	4.0
SPB-full	2	3.73	1,947	1,800	92.4	23 (0)	7.2

Table 6: Operational efficiency from the contributing primary and top-up run logs. Phase-3 attempts count materialized candidates that reached validation logging; final counts the released exactly balanced artifact; yield is final retained records divided by Phase-3 attempts. Valid Phase-3 records emitted before final category replacement/deduplication were 2,686 for Schema C and 1,946 for SPB-full. Repairs are all-phase counts with Phase-3 counts in parentheses.

the category label matches the dominant SPARQL construct. We report per-judge pass rates, inter-judge agreement, Cohen’s κ , consensus pass rate, and consensus error categories. The audit packet, raw per-judge outputs, adjudicated CSV/JSONL files, and prompt are released with the artifact bundle; the exact prompt is included in Appendix Section B.

8 Benchmark Results

The primary expanded evaluation uses Qwen/Qwen3.5-4B served through vLLM. We retain 1,800 accepted Phase-3 pairs for each schema, with exactly 200 pairs in each of the nine categories. Table 5 summarizes the strict final-artifact checks. A record is accepted only if its SPARQL parses, executes, returns a category-appropriate answer object, avoids zero-count counting artifacts, and passes category-form checks such as COUNT for counting and ranked ORDER BY/OFFSET forms for ordinal queries. These are selection-conditioned checks on the released artifacts, not a claim that every intermediate LLM proposal is correct. Both schemas satisfy the final checks without repair calls in the released accepted artifacts. The non-empty gate does not remove boolean negatives: the accepted ASK categories include 28 false answers for Schema C and 224 false answers for SPB-full.

Question diversity and structural profile. Table 7 summarizes automatic diagnostics for the

Schema	Unique Q%	Tok.	Dist-1	Dist-2	Triples	FILTER%	COUNT%	ORDER%	ASK%	OPT%
Schema C	100.0	9.8	9.8	20.3	2.9	22.7	11.6	22.2	33.3	0.0
SPB-full	100.0	10.1	10.1	21.8	2.7	21.8	24.2	22.2	33.3	9.2

Table 7: Natural-language diversity and SPARQL structural diagnostics for strict accepted Phase-3 artifacts. Dist-1/2 are corpus-level distinct unigram/bigram percentages. Per-category diagnostics are in Appendix Table 12.

Schema	Model	Pairs	Parse fail	Exec fail	Empty	SPARQL ms
Schema C	Qwen3.5-2B	450	0	0	0	4.1
Schema C	Qwen3.5-4B	450	0	0	0	16.7
Schema C	Qwen3.5-9B	450	0	0	0	5.0
SPB-full	Qwen3.5-2B	450	0	0	0	8.2
SPB-full	Qwen3.5-4B	450	0	0	0	5.8
SPB-full	Qwen3.5-9B	450	0	0	0	7.2

Table 8: Cross-model generation robustness probes at 50 pairs per category per schema.

two strict accepted artifacts. Each schema has 100% unique surface questions within the released 1,800-pair artifact. Distinct-1/2 are surface-form diagnostics rather than human naturalness scores, and should be interpreted alongside the fixed schema and category taxonomy. The SPARQL profile reflects the target taxonomy: ASK appears in boolean categories, COUNT in counting and some boolean cardinality comparisons, and ORDER BY in ranked categories. SPB-full also includes limited OPTIONAL usage in accepted comparative/difference records, while the current taxonomy does not target property paths, named graphs, or broad UNION coverage.

9 Cross-Model Robustness

To test whether the pipeline depends on a single LLM size, we rerun 50-pair-per-category probes with Qwen3.5-2B, Qwen3.5-4B, and Qwen3.5-9B on both schemas. These probes use the same strict acceptance checks as the primary runs. Table 8 shows that all six probes are balanced and produce 100% parse and execution validity with no empty answers and no repairs. This does not imply that every intermediate LLM proposal is correct; rather, it shows that reverse-query grounding plus strict validation produces reliable accepted artifacts across smaller and larger open models.

10 Semantic Audit and Benchmark Utility

The semantic audit addresses the gap between executable queries and semantically correct NL-SPARQL pairs. We evaluate a 216-pair audit packet sampled uniformly by schema and category: 2 schemas $\times$ 9 categories $\times$ 12 pairs.

Judge/decision	Pass	Pass %	Agreement	Error categories
GPT-5-mini	213/216	98.6	–	2 entity-binding, 1 intent
Grok-4.20-0309 non-reasoning	213/216	98.6	–	1 entity-binding, 2 intent
Consensus	212/216	98.1	99.1%, $\kappa = 0.662$	2 entity-binding, 2 intent

Table 9: Dual LLM-judge semantic audit over 216 stratified examples. Consensus requires both judges to mark the pair as semantically passing. Agreement and κ are computed on the overall pass/fail label; high observed agreement with moderate κ reflects the low number of failures.

The packet includes the question, SPARQL, normalized answer preview, category, and schema. Two independent LLM judges score intent/query match, entity-binding correctness, answer-type correctness, category correctness, pass/fail, and error type using the prompt in Appendix Section B. We define consensus pass as both judges marking overall_pass=yes. Consensus pass is 212/216 (98.1%). Both judges agreed on the overall pass/fail label for 214/216 records (99.1%, $\kappa = 0.662$). Consensus failures are concentrated in Schema C yes/no (10/12), SPB-full comparative (11/12), and SPB-full intersection (11/12), with adjudicated error types split between entity binding and intent mismatch.

We also evaluate whether the generated benchmark is useful for NL-to-SPARQL prompting. For each schema, we shuffle examples within each category with a fixed seed, hold out 20 examples per category (180 total), and build retrieval stores only from the remaining examples. Thus the target question and gold SPARQL are excluded from the retrieval pool. We compare a schema-only zero-shot prompt against category-retrieval prompts using $k = 3$ examples from the same schema and category. Table 10 shows large and consistent improvements from category retrieval. For Qwen3.5-2B, category-RAG raises answer F1 from 1.7 to 86.1 on Schema C and from 5.6 to 89.2 on SPB-full. For Qwen3.5-4B, category-RAG raises answer F1 from 32.8 to 91.1 on Schema C and from 18.3 to 90.0 on SPB-full. Paired bootstrap intervals over the 180 held-out examples confirm the effect: the answer-F1 gains are +84.4 points [78.9, 89.4] for Schema C 2B, +83.6 [77.8, 88.9] for SPB-full 2B, +58.3 [51.1, 65.6] for Schema C 4B, and +71.7 [65.0, 78.3] for SPB-full 4B. To separate generic example conditioning from category-matched retrieval, we add a 4B cross-category ablation that retrieves $k = 3$ examples from the same schema while excluding the target category. Cross-category RAG remains above zero-shot but below category-RAG: answer

Schema	Model	Setting	Parse%	Exec%	Answer F1	SP-F1
Schema C	Qwen3.5-2B	Zero-shot	45.0	45.0	1.7	82.5
Schema C	Qwen3.5-2B	Category-RAG	100.0	100.0	86.1	99.7
Schema C	Qwen3.5-4B	Zero-shot	75.0	73.3	32.8	89.9
Schema C	Qwen3.5-4B	Cross-cat RAG	100.0	98.3	84.2	98.5
Schema C	Qwen3.5-4B	Category-RAG	99.4	99.4	91.1	99.7
SPB-full	Qwen3.5-2B	Zero-shot	38.3	35.6	5.6	72.0
SPB-full	Qwen3.5-2B	Category-RAG	100.0	100.0	89.2	99.8
SPB-full	Qwen3.5-4B	Zero-shot	84.4	84.4	18.3	83.2
SPB-full	Qwen3.5-4B	Cross-cat RAG	99.4	99.4	78.8	97.8
SPB-full	Qwen3.5-4B	Category-RAG	100.0	100.0	90.0	99.7

Table 10: Downstream utility evaluation on 180 held-out examples per schema/model/setting. Category-RAG uses $k = 3$ generated examples from the same schema and target category, while Cross-cat RAG uses same-schema examples from non-target categories; held-out target records are excluded from every retrieval index.

F1 is 84.2 versus 91.1 on Schema C, and 78.8 versus 90.0 on SPB-full. These results support the benchmark’s practical utility as a source of schema- and category-matched examples, not only as a static evaluation set. The split prevents direct target leakage; the experiment measures the intended within-schema, category-matched retrieval setting, and Appendix Table 13 reports per-category 4B deltas.

11 Discussion and Operational Lessons

The results address the main empirical risks for schema-specific benchmark construction: scale, behavior beyond one compact graph, LLM dependence, and semantic validation. The primary benchmark contains 3,600 accepted Phase-3 pairs; the second schema uses a much larger public RDF graph with a broader type/predicate profile while holding the task taxonomy fixed; and cross-model probes show that strict accepted artifacts are not tied to one Qwen 3.5 model size. This should be read as graph-scale and schema-profile evidence, not as exhaustive use of every construct in SPB. The downstream utility experiment adds a separate usefulness claim: generated category-matched examples materially improve held-out NL-to-SPARQL prompting, especially for smaller models, and the cross-category ablation shows that category matching adds value beyond generic same-schema examples.

The most important operational lesson is that SPARQL endpoint design matters as much as model choice when scaling benchmark generation. The initial large-graph runs exposed endpoint bottlenecks from randomized reverse queries and many tiny metadata lookups. Replacing ORDER BY RAND() with deterministic binding banks, shuffling locally, and batching label/type retrieval re-

duced GraphDB load enough to complete balanced 200-per-category runs. This supports the pipeline’s methodological claim: execution grounding is practical when reverse-query work is treated as a planned data-building stage rather than an on-line sampling side effect.

The dual LLM audit gives a sampled semantic check on top of parse and execution validity. The adjudicated consensus failures are two entity-binding issues and two intent mismatches. Two involve judge disagreement over role-like key-person labels, and two are full consensus failures involving same-location inequality and role binding. These are exactly the kinds of errors that execution-only checks can miss, making semantic judging useful both as an evaluation result and as an automated filtering or top-up signal for future benchmark refreshes.

12 Benchmark Usage

PIPE-RDF produces benchmark artifacts that can be used to evaluate NL-to-SPARQL systems on a target RDF schema. The generated pairs support execution accuracy, answer F1, SP-F1, and structural similarity evaluation. The same seed banks can also be reused as few-shot retrieval examples or synthetic supervision for schema-specific fine-tuning. The method is especially useful when public KGQA datasets do not match the target ontology, but the contribution does not depend on demonstrating adoption in any particular deployment context.

13 Conclusion

PIPE-RDF provides a practical pipeline for generating balanced, schema-specific NL-SPARQL benchmarks. Across two RDF schemas, the primary Qwen3.5-4B runs produce 3,600 accepted Phase-3 pairs with exact category balance, no parse failures, no execution failures, no empty answers, and no repairs in the released accepted artifacts. Cross-model probes with Qwen3.5-2B, 4B, and 9B show the same strict accepted-artifact profile at smaller scale, and the downstream utility evaluation shows that generated category-matched examples substantially improve NL-to-SPARQL prompting. By combining reverse querying, category-aware retrieval, binding banks, and execution-based validation, PIPE-RDF supports reproducible benchmark construction for RDF question answering.

14 Limitations

Both evaluated schemas are public. Schema C is a controlled public company-location slice and SPB-full is a public benchmark graph; we do not claim evidence from private or proprietary deployments. The paper’s claim is about schema-specific benchmark construction, not about the prevalence or adoption of RDF/SPARQL in any particular sector.

The current query taxonomy covers ASK, COUNT, FILTER, ORDER/LIMIT/OFFSET, VALUES, DISTINCT, joins, and limited OPTIONAL use in accepted SPB records. It does not yet target property paths, named graphs, language-tag filters, broad UNION usage, or full NOT EXISTS/subquery-centered categories. These are natural extensions of the template, binding-bank, and category-validator interface rather than changes to the core execution-grounded design. The present runs use English labels and questions; multilingual extension would require language-tag-aware label retrieval, multilingual prompt templates, and language-aware semantic judging.

The downstream utility evaluation is prompt-based rather than a full comparison of fine-tuned semantic parsers. It shows that generated examples improve retrieval-augmented prompting, but future work should evaluate fine-tuning, grammar-constrained decoding, and additional RDF schemas with different ontology styles. The Qwen3.5-9B probe required a shorter 8K context on one GPU because the 16K configuration exceeded available memory; all probe prompts fit within that context, but larger-context deployments should report serving constraints explicitly.

15 Ethical Considerations

PIPE-RDF is intended for benchmark construction over RDF graphs. When used with private graphs, deployments should follow data governance policies, minimize exposure of sensitive identifiers in prompts, and audit logs for leakage. The public schemas in this study are derived from public RDF data, and usage follows their licensing terms.

References

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child,

Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. [Language models are few-shot learners](#). *Preprint*, arXiv:2005.14165.

Shulin Cao, Jiaxin Shi, Liangming Pan, Lunyiu Nie, Yutong Xiang, Lei Hou, Juanzi Li, Bin He, and Hanwang Zhang. 2022. [KQA pro: A dataset with explicit compositional programs for complex question answering over knowledge base](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6101–6119, Dublin, Ireland. Association for Computational Linguistics.

Mohnish Dubey, Debayan Banerjee, Abdelrahman Abdelkawi, and Jens Lehmann. 2019. [LC-QuAD 2.0: A large dataset for complex question answering over wikidata and dbpedia](#). In *The Semantic Web - ISWC 2019*, Lecture Notes in Computer Science, pages 69–78. Springer.

Jonathan Fürst, Cedric Osei-Akoto, Sithursan Sivasubramaniam, Kurt Stockinger, and Yi Zhang. 2024. [SM3-text-to-query: Synthetic multi-model medical text-to-query benchmark](#). In *Advances in Neural Information Processing Systems 37*, pages 88627–88663. Neural Information Processing Systems Foundation, Inc. (NeurIPS).

Daniel Keysers, Nathanael Schärli, Nathan Scales, Hylke Buisman, Daniel Furrer, Sergii Kashubin, Nikola Momchev, Danila Sinopalnikov, Lukasz Stafiniak, Tibor Tihon, Dmitry Tsarkov, Xiao Wang, Marc van Zee, and Olivier Bousquet. 2020. [Measuring compositional generalization: A comprehensive method on realistic data](#). *Preprint*, arXiv:1912.09713.

Catherine Kosten, Philippe Cudré-Mauroux, and Kurt Stockinger. 2023. [Spider4sparql: A complex benchmark for evaluating knowledge graph question answering systems](#). In *IEEE International Conference on Big Data, BigData 2023*, pages 5272–5281. IEEE.

Jaebok Lee and Hyeonjeong Shin. 2025. [SPARKLE: Enhancing SPARQL generation with direct KG integration in decoding](#). *Expert Systems with Applications*, 289:128263.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2021. [Retrieval-augmented generation for knowledge-intensive nlp tasks](#). *Preprint*, arXiv:2005.11401.

Linked Data Benchmark Council. 2026. [LDBC semantic publishing benchmark \(SPB\)](#). <https://ldbcouncil.org/benchmarks/spb>. Accessed 2026-05-20.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. [G-eval](#):

768	NLG evaluation using gpt-4 with better human alignment.	823
769	In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing</i> ,	824
770	pages 2511–2522, Singapore. Association for Computational Linguistics.	825
771		826
772		
773	Reham Omar, Abdelghny Orogat, Ibrahim Abdelaziz,	827
774	Omij Mangukiya, Panos Kalnis, and Essam Mansour.	828
775	2025. <i>Chatty-KG: A multi-agent AI system for on-demand conversational question answering over knowledge graphs</i> . <i>Preprint</i> , arXiv:2511.20940.	829
776		830
777		831
778	Xueli Pan, Victor de Boer, and Jacco van Ossenbruggen.	832
779	2025. <i>FIRESPARQL: A LLM-based framework for SPARQL query generation over scholarly knowledge graphs</i> . <i>Preprint</i> , arXiv:2508.10467.	833
780		834
781		
782	Aleksandr Perevalov, Dennis Diefenbach, Ricardo Usbeck,	835
783	and Andreas Both. 2022. <i>Qald-9-plus: A multilingual dataset for question answering over dbpedia and wikidata translated by native speakers</i> . <i>Preprint</i> ,	836
784	arXiv:2202.00120.	837
785		838
786		839
787		840
788	Shuo Ren, Daya Guo, Shuai Lu, Long Zhou, Shujie Liu,	841
789	Duyu Tang, Neel Sundaresan, Ming Zhou, Ambrosio Blanco,	842
790	and Shuai Ma. 2020. <i>Codebleu: a method for automatic evaluation of code synthesis</i> . <i>Preprint</i> ,	843
791	arXiv:2009.10297.	844
792		845
793		846
794	Md Rashad Al Hasan Rony, Uttam Kumar, Roman Teucher,	847
795	Liubov Kovriguina, and Jens Lehmann. 2022. <i>Sgpt: A generative approach for sparql query generation from natural language questions</i> . <i>IEEE Access</i> , 10:70712–70723.	848
796		
797	Muhammad Saleem, Qaiser Mehmood, and Axel-Cyrille Ngonga Ngomo. 2015. <i>FEASIBLE: A feature-based SPARQL benchmark generation framework</i> . In <i>The Semantic Web - ISWC 2015</i> , Lecture Notes in Computer Science, pages 52–69. Springer.	849
798		850
799		851
800		852
801		853
802	Michael Schmidt, Thomas Hornung, Georg Lausen, and Christoph Pinkel. 2008. <i>SP2Bench: A SPARQL performance benchmark</i> . <i>Preprint</i> , arXiv:0806.4627.	854
803		855
804		
805	Torsten Scholak, Nathan Schucher, and Dzmitry Bahdanau. 2021. <i>Picard: Parsing incrementally for constrained auto-regressive decoding from language models</i> . In <i>Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing</i> , pages 9895–9901. Association for Computational Linguistics.	856
806		857
807		858
808		859
809		
810		860
811		861
812	Priyanka Sen, Alham Fikri Aji, and Amir Saffari. 2022. <i>Mintaka: A complex, natural, and multilingual dataset for end-to-end question answering</i> . <i>Preprint</i> ,	862
813	arXiv:2210.01613.	863
814		864
815		865
816	Tommaso Soru, Edgard Marx, Diego Moussallem, Gustavo Publio, André Valdestilhas, Diego Esteves, and	866
817	Ciro Baron Neto. 2017. <i>SPARQL as a foreign language</i> . <i>Preprint</i> , arXiv:1708.07624.	867
818		
819		
820	Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier	868
821	Martinet, Marie-Anne Lachaux, Timothée Lacroix,	869
822	Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal	870
	Azhar, Aurelien Rodriguez, Armand Joulin, Edouard	871
	Grave, and Guillaume Lample. 2023. <i>Llama: Open and efficient foundation language models</i> . <i>Preprint</i> , arXiv:2302.13971.	
	Priyansh Trivedi, Gaurav Maheshwari, Mohnish Dubey,	
	and Jens Lehmann. 2017. <i>LC-QuAD: A Corpus for Complex Question Answering over Knowledge Graphs</i> , pages 210–218. Springer International Publishing.	
	Sebastian Walter and Hannah Bast. 2025. <i>GRASP: Generic reasoning and SPARQL generation across knowledge graphs</i> . <i>Preprint</i> , arXiv:2507.08107.	
	Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. <i>Judging LLM-as-a-judge with MT-bench and chatbot arena</i> . <i>Preprint</i> , arXiv:2306.05685.	
	Xiwen Zheng, Arun Kumar, and Amarnath Gupta. 2024. <i>Generating cross-model analytics workloads using llms</i> . In <i>Proceedings of the 33rd ACM International Conference on Information and Knowledge Management</i> , pages 4303–4307. Association for Computing Machinery.	
	Appendix	
	A Reproducibility and Artifact Manifest	
	• Software/models: GraphDB 10.6.3 (Docker), vLLM service, Qwen 3.5 chat models, embedding model BAAI/bge-m3.	
	• Endpoints/repositories: spb_company_mini_slice for Schema C and spb_1m for SPB-full, both queried with inference disabled (infer=false).	
	• Primary configs: the ARR config directory contains arr_schema_c_200, arr_spb_full_200, and matching arr_cross_model YAMLs.	
	• Randomness control: sampling utilities use seed 42; LLM generation is stochastic, so released JSONL/CSV artifacts are the exact evaluation objects.	
	• Command template: run the pipeline entry point with an ARR YAML config; utility and audit entry points are released under scripts/.	
	• Artifacts: run directories contain logs, phase-wise JSONL/CSV files, binding banks, summaries, and analysis outputs under artifacts/runs/<run_id>/.	

- **Hardware/deployment:** GraphDB in Docker plus vLLM model serving on GPU resources; endpoint and model-serving scripts are included under scripts/ds_serv6_*.sh.
- **Semantic audit:** the 216-record packet, raw judge outputs, adjudicated files, and summary are released under the LLM audit artifact group; the logged judge identifiers are gpt-5-mini and grok-4.20-0309 non-reasoning.

Evidence	Artifact ID	Summary file
Schema C primary	RUN_SCHEMA_C_4B	benchmark_phase3_balanced_summary.json
SPB-full primary	RUN_SPB_4B	benchmark_phase3_balanced_summary.json
Cross-model probes	RUN_CROSS_MODEL	run_summary.json
LLM semantic audit	AUDIT_LLM_FINAL_V2	llm_semantic_audit_summary.json
NL structural diagnostics	DIAG_STRUCTURAL	linguistic_structural_diagnostics.json
Operational efficiency	DIAG_EFFICIENCY	operational_efficiency_summary.json
Downstream category stats	DIAG_UTILITY_CATEGORY	downstream_by_category_summary.json
Schema profiles	PROFILE_SCHEMA	summary.json

Table 11: Released artifact groups used to populate the main tables. Full directory names are preserved in the artifact bundle and README; short IDs avoid brittle path wrapping in the paper.

B LLM Semantic Judge Prompt

The semantic audit uses the following prompt for both LLM judges; fields in braces are filled from each audit record.

System prompt:

You are an expert evaluator of NL-SPARQL benchmark quality.

Your task is to judge whether a natural-language question and its SPARQL query are semantically aligned for the stated RDF schema and query category. Focus on meaning, bindings, answer type, and category construct. Do not reward a query merely because it parses or executes. Use the answer preview only as supporting evidence; the primary decision is whether the SPARQL query correctly operationalizes the question on the given schema.

Return JSON only. Do not include markdown, prose, or code fences.

User prompt template:

Evaluate this PIPE-RDF NL-SPARQL benchmark record.

Audit ID: {audit_id}
Schema: {schema}
Category: {category}

Natural-language question:
{question}

SPARQL query:
```sparql  
{sparql}  
```

Answer preview JSON:
{answers}

Pipeline checks:
- parse_valid: {source_parse_valid}
- execution_success: {source_exec_success}
- answer_count: {source_answer_count}

Rubric:

1. intent_query_match: yes iff the query asks for the same relation, comparison, count, set operation, ranking, or yes/no condition as the question.
2. entity_binding_correct: yes iff concrete entities or literals in the question are correctly represented by VALUES clauses, IRIs, filters, or equivalent query constraints. If the question has no concrete entity, mark yes only if the query does not introduce a contradictory concrete entity.
3. answer_type_correct: yes iff the selected output includes the answer type requested by the question, such as entity, literal, boolean, count, date/year, or ranked item. Supporting columns are acceptable: for example, a superlative query may return both the winning entity and the numeric value used for ranking.
4. category_construct_correct: yes iff the SPARQL construct matches the intended category. Examples: ASK for yes/no and boolean comparative/difference questions; COUNT(DISTINCT ...) for counting; ORDER BY with LIMIT/OFFSET for superlative or ordinal; connected multi-triple joins through variables for multi-hop; conjunction for intersection; exclusion or inequality for difference. For multi-hop, direct schema predicates are valid when they are the schema path that operationalizes the question; do not require an unrelated ontology path.
5. overall_pass: yes iff all required semantic checks pass and there is no material issue.

When overall_pass is no, choose exactly one primary error_type from:
intent_mismatch, entity_binding, answer_type, category_construct, syntax_or_execution, unsupported_or_ambiguous_question, other.
When overall_pass is yes, use error_type none.

Judgment conventions:

- Treat canonical IRIs and aliases as correct when they denote the same entity named in the question.
- Do not mark answer_type_correct no only because the query returns an auxiliary label, year, count, or numeric value in addition to the requested entity.
- Do not mark category_construct_correct no only because the query uses the schema's direct predicate names; judge whether the variables and joins express the intended category under the given schema.

Return exactly this JSON object:

```
{
  "intent_query_match": "yes or no",
  "entity_binding_correct": "yes or no",
  "answer_type_correct": "yes or no",
  "category_construct_correct": "yes or no",
  "overall_pass": "yes or no",
  "error_type": "none or one allowed error type",
  "confidence": 0.0,
  "notes": "one short sentence"
}
```

C Schema Figure

Schema C: Company-Location Mini-Slice




Figure 2: Schema C mini-slice used in this study. **dbo:Company** connects to typed **gn:Feature** locations and **foaf:Person** key people; industry is represented as a labeled IRI-valued attribute and may not have an explicit type in the slice.

Category	Unique Q%	Tok.	Dist-2	Triples	FILTER%	COUNT%	ORDER%	ASK%	OPT%
Comparative	100.0	12.5	23.7	3.7	98.3	26.0	0.0	100.0	13.0
Counting	91.5	6.9	25.4	2.0	0.3	100.0	0.0	0.0	0.0
Difference	100.0	11.8	26.2	3.4	100.0	32.0	0.0	100.0	28.3
Generic	100.0	6.5	49.0	2.0	2.0	0.0	0.0	0.0	0.0
Intersection	100.0	12.3	19.9	3.0	0.0	1.5	0.0	0.0	0.0
Multi-hop	100.0	11.8	15.3	3.0	0.0	0.5	0.0	0.0	0.0
Ordinal	79.0	9.5	7.8	3.0	0.0	0.0	100.0	0.0	0.0
Superlative	90.0	8.8	20.4	3.0	0.0	0.5	100.0	0.0	0.0
Yes/No	100.0	9.4	49.3	2.1	0.0	0.5	0.0	100.0	0.0

Table 12: Combined per-category natural-language and SPARQL structural diagnostics for the two strict primary artifacts. Unique-question rates are computed after pooling both schemas; per-schema uniqueness is 100% for every category, while pooled counting/ordinal/superlative rates are lower because the same high-level pattern can occur in both schemas.

Schema	Category	Answer F1 Δ	95% CI	Sign-test p
Schema C	Comparative	+30.0	[10.0, 50.0]	0.031
Schema C	Counting	+95.0	[85.0, 100.0]	< 0.001
Schema C	Difference	+15.0	[0.0, 35.0]	0.250
Schema C	Generic	+15.0	[0.0, 35.0]	0.250
Schema C	Intersection	+95.0	[85.0, 100.0]	< 0.001
Schema C	Multi-hop	+75.0	[55.0, 90.0]	< 0.001
Schema C	Ordinal	+95.0	[85.0, 100.0]	< 0.001
Schema C	Superlative	+90.0	[75.0, 100.0]	< 0.001
Schema C	Yes/No	+15.0	[0.0, 30.0]	0.250
SPB-full	Comparative	+50.0	[30.0, 70.0]	0.002
SPB-full	Counting	+85.0	[65.0, 100.0]	< 0.001
SPB-full	Difference	+55.0	[35.0, 75.0]	< 0.001
SPB-full	Generic	+15.0	[0.0, 30.0]	0.250
SPB-full	Intersection	+80.0	[60.0, 95.0]	< 0.001
SPB-full	Multi-hop	+70.0	[50.0, 90.0]	< 0.001
SPB-full	Ordinal	+100.0	[100.0, 100.0]	< 0.001
SPB-full	Superlative	+90.0	[75.0, 100.0]	< 0.001
SPB-full	Yes/No	+100.0	[100.0, 100.0]	< 0.001

Table 13: Per-category category-RAG minus zero-shot answer-F1 deltas for Qwen3.5-4B. Each category has 20 paired held-out examples; confidence intervals are paired bootstrap intervals.

D Example NL-SPARQL Pairs

We include representative Phase-3 examples covering entity lookup, aggregation, boolean comparison, ranking, and conjunction. Prefix declarations are omitted for brevity; the artifact bundle contains full examples for all nine categories.

Generic. Q: Who is a key person at (Company: Volkswagen)?

```
SELECT ?name
WHERE {
  VALUES ?company {
    ↪ <http://dbpedia.org/resource/Volkswagen> }
  ?company rdf:type dbo:Company .
  ?company dbo:keyPerson ?person .
  ?person rdf:type foaf:Person .
  ?person foaf:name ?name .
}
LIMIT 5
```

Counting. Q: How many companies are located in (Location: Menlo Park, California)?

```
SELECT (COUNT(DISTINCT ?company) AS ?count)
WHERE {
  VALUES ?location { <http://dbpedia.org/resour_
    ↪ ce/Menlo_Park,_California> }
  ?company a dbo:Company .
  ?company dbo:location ?location .
  ?location a gn:Feature .
}
```

Comparative. Q: Do (Company: Airbus) and (Company: FileMaker Inc.) have different numbers of employees?

```
ASK {
  VALUES ?company1 {
    ↪ <http://dbpedia.org/resource/Airbus> }
  VALUES ?company2 { <http://dbpedia.org/resour_
    ↪ ce/FileMaker_Inc.> }
  ?company1 a dbo:Company .
  ?company2 a dbo:Company .
  ?company1 dbo:numberOfEmployees ?n1 .
  ?company2 dbo:numberOfEmployees ?n2 .
  FILTER (xsd:integer(?n1) != xsd:integer(?n2))
}
```

Ordinal. Q: Which company was founded second earliest?

```
SELECT ?company ?year
WHERE {
  ?company a dbo:Company .
  ?company dbo:foundingYear ?year .
}
ORDER BY ASC(xsd:integer(?year)) ?company
OFFSET 1
LIMIT 1
```

Intersection. Q: Which companies are located in (Location: California) and are in the (Industry: Software)?

```
SELECT DISTINCT ?company
WHERE {
  ?company a dbo:Company .
  ?company dbo:location
    ↪ <http://dbpedia.org/resource/California> .
  ?company dbo:industry
    ↪ <http://dbpedia.org/resource/Software> .
}
LIMIT 5
```